


# Optimizations which made Python 3.6 faster than Python 3.5



Pycon US 2017, Portland, OR



redhat®

Victor Stinner  
vstinner@redhat.com

# Agenda

---



- (1) Benchmarks
- (2) Benchmarks results
- (3) Python 3.5 optimizations
- (4) Python 3.6 optimizations
- (5) Python 3.7 optimizations

# Agenda

---



## (1) Benchmarks



# Unstable benchmarks

---



- March 2016, no developer trusted the Python benchmark suite
- Many benchmarks were unstable
- It wasn't possible to decide if an optimization makes CPython faster or not...

# New perf module

---



- Calibrate the number of loops
- Spawn **20 processes** sequentially, **3** values per process, total: **60** values
- Compute **average** (mean) and **standard deviation**

# performance project



- Benchmarks rewritten using perf: new project **performance** on GitHub
- <http://speed.python.org> now runs performance
- CPython is now compiled with Link Time Optimization (**LTO**) and Profile Guided Optimization (**PGO**)




# Linux and CPUs

---



- **sudo python3 -m perf system tune**
- Use fixed CPU frequency
- Disable Intel Turbo Boost
- If CPU isolation is enabled, Linux kernel options `isolcpus` and `rcu_nocbs`, use CPU pinning
- CPU isolation helps a lot to reduce operation system jitter

# Spot perf regression




python\_startup: 20 ms => 27 ms, fix: 17 ms





# Timeline



April, 2014 – May, 2017: 3 years



# Agenda

---




## (2) Benchmarks results

# 3.6 faster than 3.5



Time normalized to lto-pgo latest in branch '3.5'




Results normalized to Python 3.5

lower = faster






# 3.6 faster than 2.7



Results normalized to Python 2.7  
lower = faster




# telco: 3.6 vs 2.7



Python 3.6 is **40x** faster than Python 2.7  
(decimal module rewritten in C  
by Stefan Krah in Python 3.3)

# 3.7 faster than 3.6



Results normalized to Python 3.6

lower = faster





# Agenda

---



## (3) Python 3.5 optimizations

# lru\_cache()



- Matt Joiner, Alexey Kachayev and Serhiy Storchaka reimplemented `functools.lru_cache()` in C
- sympy: **20% faster**
- scimark\_lu: **5% faster**
- Tricky C code, hard to get it right: 3 years  $\frac{1}{2}$  to close the bpo-14373

# OrderedDict

---



- Eric Snow reimplemented collections.OrderedDict in C
- html5lib: **20% faster**
- Reuse C implementation of dict
- Again, tricky C code: 2 years ½ to close the bpo-16991



# Agenda

---



## (4) Python 3.6 optimizations

# PyMem\_Malloc()



- Victor Stinner changed `PyMem_Malloc()` to use Python fast memory allocator
- Many benchmarks: **5% - 22% faster**
- Check if the GIL is held in debug hooks
- Only numpy misused the API (fixed)
- **PYTHONMALLOC=debug** now available in release builds to detect memory corruptions, bpo-26249

# ElementTree parse

---



- Serhiy Storchaka optimized `ElementTree.iterparse()`
- **2x faster**
- Follow-up of Brett Canon's Pycon Canada 2015 keynote :-)
- bpo-25638



# PGO uses tests

---



- Brett Canon modified the Profile Guided Optimization (PGO)
- The Python test suite is now used, rather than pidigits, to guide the compiler
- Many benchmarks: 5% – 27% faster
- bpo-24915

# Wordcode

---



- Demur Rumed and Serhiy Storchaka modified the bytecode to always use 2 bytes opcodes
- Before: 1 (no arg) or 3 bytes (with arg)
- Removed an if from ceval.c hotcode for **better CPU branch prediction**:  

```
if (HAS_ARG(opcode))  
    oparg = NEXTARG();
```
- bpo-26647

# FASTCALL

---



- Victor Stinner wrote a new C API to avoid the creation of temporary tuples to pass function arguments
- Many microbenchmarks: **12% – 50% faster**
- `obj[0], getattr(obj, "attr"), {1: 2}.get(1), list.count(0), str.replace("a", "b"), ...`
- Avoid **20 ns** per modified function call



# Unicode codecs

---



- Victor Stinner optimized ASCII and UTF-8 codecs for ignore, replace, surrogateescape and surrogatepass error handlers
- UTF-8: decoder **15x faster**, encoder **75x faster**
- ASCII: decoder **60x faster**, encoder **3x faster**



# bytes % args



- PEP 461 added back bytes % args to Python 3.5
- Victor Stinner wrote a new `_PyBytesWriter` API to optimize functions creating bytes and bytearray strings
- bytes % args: **2x faster**
- bytes.fromhex(): **3x faster**

# Globbing

---



- Serhiy Storchaka optimized `glob.glob()`, `glob.iglob()` and pathlib globbing using `os.scandir()` (new in Python 3.5)
- `glob`: **3x - 6x faster**
- Pathlib `glob`: **1.5x - 4x faster**
- Avoid one `stat()` per directory entry
- bpo-25596, bpo-26032

# asyncio

---



- Yury Selivanov and Naoki INADA reimplemented asyncio Future and Task classes in C
- Asyncio programs: **30% faster**
- bpo-26081, bpo-28544



# Agenda

---



## (5) Python 3.7 optimizations



# Method calls

---



- Yury Selivanov and Naoki INADA added LOAD\_METHOD and CALL\_METHOD opcodes
- Methods calls: **10% - 20% faster**
- Idea coming from PyPy, bpo-26110


# Future optimizations

---



- More optimizations are coming in Python 3.7...
- Stay tuned!

# 3.7 slower than 2.7 :-)



Results normalized to Python 2.7

higher = slower



# Questions?

---



<http://speed.python.org/>

<http://faster-cpython.readthedocs.io/>