


Optimizations which made Python 3.6 faster than Python 3.5



Pycon US 2017, Portland, OR



redhat®

Victor Stinner
vstinner@redhat.com

Agenda



- (1) Benchmarks
- (2) Benchmarks results
- (3) Python 3.5 optimizations
- (4) Python 3.6 optimizations
- (5) Python 3.7 optimizations

Agenda



(1) Benchmarks

Unstable benchmarks



- March 2016, no developer trusted the Python benchmark suite
- Many benchmarks were unstable
- It wasn't possible to decide if an optimization makes CPython faster or not...

New perf module



- Calibrate the number of loops, at least 100 ms
- Spawn **20 processes** sequentially
- Each process runs the benchmark once to warmup and then run the benchmark 3 times
- Total: **60** values. Compute **average** (mean) and **standard deviation**

New perf module



- Write results into JSON
- Compare results
- Statistics to check for outliers, min/max, percentiles, etc.
- Render an histogram
- Command to dump all values
- ...

performance project




- Benchmarks rewritten using perf: new project **performance** on GitHub
- <http://speed.python.org> now runs performance
- CPython is now compiled with Link Time Optimization (**LTO**) and Profile Guided Optimization (**PGO**)

Linux and CPUs



- **sudo python3 -m perf system tune**
- Use fixed CPU frequency, disable Intel Turbo Boost
- If CPU isolation is enabled, Linux kernel options `isolcpus` and `rcu_nocbs`, use CPU pinning
- CPU isolation helps a lot to reduce operation system jitter


Spot perf regression



python_startup: 20 ms => 27 ms, fix: 17 ms



Timeline



April, 2014 – May, 2017: 3 years



Agenda




(2) Benchmarks results

3.6 faster than 3.5



Time normalized to lto-pgo latest in branch '3.5'




Results normalized to Python 3.5

lower = faster




3.6 faster than 2.7



Results normalized to Python 2.7
lower = faster




telco: 3.6 vs 2.7



Python 3.6 is **40x** faster than Python 2.7
decimal module rewritten in C

3.7 faster than 3.6



Results normalized to Python 3.6

lower = faster



Agenda



(3) Python 3.5 optimizations

lru_cache()



- Matt Joiner, Alexey Kachayev and Serhiy Storchaka reimplemented `functools.lru_cache()` in C
- sympy: **20% faster**
- scimark_lu: **5% faster**
- Tricky C code, hard to get it right: 3 years $\frac{1}{2}$ to close the bpo-14373

OrderedDict



- Eric Snow reimplemented collections.OrderedDict in C
- html5lib: **20% faster**
- Reuse C implementation of dict
- Again, tricky C code: 2 years ½ to close the bpo-16991

Agenda



(4) Python 3.6 optimizations

PyMem_Malloc()



- Victor Stinner changed `PyMem_Malloc()` to use Python fast memory allocator
- Many benchmarks: **5% - 22% faster**
- Check if the GIL is held in debug hooks
- Only numpy misused the API (fixed)
- **PYTHONMALLOC=debug** now available in release builds to detect memory corruptions, bpo-26249

ElementTree parse



- Serhiy Storchaka optimized `ElementTree.iterparse()`
- **2x faster**
- Follow-up of Brett Canon's Pycon Canada keynote :-)
- bpo-25638

PGO uses tests



- Brett Canon modified the Profile Guided Optimization (PGO)
- The Python test suite is now used, rather than pidigits, to guide the compiler
- Many benchmarks: 5% – 27% faster
- bpo-24915

Wordcode



- Demur Rumed and Serhiy Storchaka modified the bytecode to always use 2 bytes opcodes
- Before: 1 (no arg) or 3 bytes (with arg)
- Removed an if from ceval.c hotcode for **better CPU branch prediction**:

```
if (HAS_ARG(opcode))  
    oparg = NEXTARG();
```
- bpo-26647

FASTCALL



- Victor Stinner wrote a new C API to avoid the creation of temporary tuples to pass function arguments
- Many microbenchmarks: **12% – 50% faster**
- `obj[0], getattr(obj, "attr"), {1: 2}.get(1), list.count(0), str.replace("a", "b"), ...`
- Avoid **20 ns** per modified function call

Unicode codecs



- Victor Stinner optimized ASCII and UTF-8 codecs for ignore, replace, surrogateescape and surrogatepass error handlers
- UTF-8: decoder **15x faster**, encoder **75x faster**
- ASCII: decoder **60x faster**, encoder **3x faster**

bytes % args



- PEP 461 added back bytes % args to Python 3.5
- Victor Stinner wrote a new `_PyBytesWriter` API to optimize functions creating bytes and bytearray strings
- bytes % args: **2x faster**
- bytes.fromhex(): **3x faster**

Globbing



- Serhiy Storchaka optimized `glob.glob()`, `glob.iglob()` and pathlib globbing using `os.scandir()` (new in Python 3.5)
- `glob`: **3x - 6x faster**
- Pathlib `glob`: **1.5x - 4x faster**
- Avoid one `stat()` per directory entry
- bpo-25596, bpo-26032

asyncio



- Yury Selivanov and Naoki INADA reimplemented asyncio Future and Task classes in C
- Asyncio programs: **1.3x faster**
- bpo-26081, bpo-28544

Agenda



(5) Python 3.7 optimizations

Method calls




- Yury Selivanov and Naoki INADA added LOAD_METHOD and CALL_METHOD opcodes
- Methods calls: **10% - 20% faster**
- Idea coming from PyPy, bpo-26110

Future optimizations



- More optimizations are coming in Python 3.7...
- Stay tuned!

3.7 slower than 2.7 :-)



Results normalized to Python 2.7

higher = slower



Questions?



<http://speed.python.org/>

<http://faster-cpython.readthedocs.io/>