


Optimizations which made Python 3.6 faster than Python 3.5



Pycon US 2017, Portland, OR



redhat®

Victor Stinner
vstinner@redhat.com

Agenda



- (1) Benchmarks
- (2) Benchmarks results
- (3) Python 3.5 optimizations
- (4) Python 3.6 optimizations
- (5) Python 3.7 optimizations

Agenda



(1) Benchmarks

(1) Unstable benchmarks



- March 2016, no developer trusted the Python benchmark suite
- Many benchmarks were unstable
- It wasn't possible to decide if an optimization makes CPython faster or not...

(1) perf module



- New **perf** project: spawn **20 processes**, each runs the benchmark once to warmup and then run it 3 times
- Total of 60 timings: compute **average** (mean) and **standard deviation**
- Statistics to check for outliers, min/max, percentiles, etc.

(1) performance project


- Benchmarks rewritten using perf: new project **performance** on GitHub
- <http://speed.python.org> now runs performance
- CPython is now compiled with Link Time Optimization (**LTO**) and Profile Guided Optimization (**PGO**)

(1) Linux and CPUs



- **python3 -m perf system tune**
- Use fixed CPU frequency, disable Intel Turbo Boost
- Use CPU isolation (Linux `isolcpus` and `rcu_nocbs`), pin benchmark process and IRQs on CPUs
- Reduce Linux perf sampling rate
- Laptop: check that power cable is plugged


(1) Spot perf regression



python_startup: 20 ms => 27 ms => 17 ms



(1) Timeline



April, 2014 – May, 2017: 3 years




Agenda



(2) Benchmarks results

(2) 3.6 faster than 3.5




Results normalized to Python 3.5

lower = faster




(2) 3.6 faster than 2.7




Results normalized to Python 2.7
lower = faster

(2) telco: 3.6 vs 2.7



Python 3.6 is **40x** faster than Python 2.7
decimal module rewritten in C

(2) 3.7 faster than 3.6



Results normalized to Python 3.6

lower = faster



Agenda



(3) Python 3.5 optimizations

(3) `C lru_cache()`



- Matt Joiner, Alexey Kachayev and Serhiy Storchaka reimplemented `functools.lru_cache()` in C
- sympy: **1.2x faster**
- scimark_lu: **1.06x faster**
- bpo-14373

(3) C OrderedDict



- Eric Snow reimplemented collections.OrderedDict in C
- html5lib: **1.2x faster**
- bpo-16991

Agenda



(4) Python 3.6 optimizations

(4) PyMem_Malloc()



- Victor Stinner modified `PyMem_Malloc()` to reuse Python fast memory allocator
- Many benchmarks: **1.05x - 1.36x faster**
- **PYTHONMALLOC=debug** now available in release builds to detect memory corruptions
- bpo-26249

(4) ElementTree parse



- Serhiy Storchaka optimized ElementTree.iterparse()
- **2x faster**
- bpo-25638

(4) PGO uses tests



- Brett Canon modified the build system to guide the compiler using the Python test suite rather than pidigits for the Profile Guided Optimization (PGO)
- Many benchmarks: 1.05x - 1.28x faster
- bpo-24915

(4) Wordcode



- Demur Rumed and Serhiy Storchaka modified the bytecode to always use 16-bit instructions:
8 bit op, 8 bit arg (arg=0 if unused)
- Removed an if from ceval.c hotcode for **better CPU branch prediction**:

```
if (HAS_ARG(opcode))  
    oparg = NEXTARG();
```
- bpo-26647

(4) FASTCALL



- Victor Stinner wrote a new C API to avoid the creation of temporary tuples to pass function arguments
- Many microbenchmarks: **1.13x - 1.92x faster**
- `obj[0], getattr(obj, "attr"), {1: 2}.get(1), ...`
- Avoid **20 ns** per modified function call

(4) Unicode codecs



- Victor Stinner optimized ASCII and UTF-8 codecs for ignore, replace, surrogateescape and surrogatepass error handlers
- UTF-8: decoder **15x faster**, encoder **75x faster**
- ASCII: decoder **60x faster**, encoder **3x faster**

(4) bytes % args



- Victor Stinner wrote a new private `_PyBytesWriter` API to optimize functions creating bytes and bytearray strings
- `bytes % args`: **2x faster**
- `bytes.fromhex()`: **3x faster**

(4) Globbing



- Serhiy Storchaka optimized `glob.glob()`, `glob.iglob()` and `pathlib` globbing
- `glob`: **3x - 6x faster**
- `Pathlib glob`: **1.5x - 4x faster**
- `bpo-25596`, `bpo-26032`

(4) C asyncio



- Yury Selivanov and Naoki INADA reimplemented asyncio Future and Task classes in C
- Asyncio programs: **1.3x faster**
- bpo-26081, bpo-28544

Agenda



(5) Python 3.7 optimizations

(5) Method calls



- Yury Selivanov and Naoki INADA added LOAD_METHOD and CALL_METHOD opcodes
- Methods calls: **1.2x faster**
- Idea coming from PyPy, bpo-26110

(5) Future optimizations



- More optimizations are coming in Python 3.7...
- Stay tuned!

Questions?



<http://speed.python.org/>

<http://faster-cpython.readthedocs.io/>